HaltIQ optional Jev comparison

OPTIONAL HOSTED COMPANION - SEPARATE FROM PRIMARY

Evidence checkpoint: 2026-10-18T20:51:17.106638-05:00

Observed outcome and matched protocol

The optional hosted Jev arm showed better observed first-decision correctness discrimination on this small panel, while answer quality stayed unchanged: every listed policy scored 10/20 exact matches (50%) and token F1 0.540. No EM repair or regression occurred. These observations establish neither general superiority nor population equivalence.

This companion is separate from the primary calibrated Strands comparison. It reuses the identical 20 frozen public-validation HotpotQA questions, order, full supplied evidence, references and generic ten-draft trajectories (initial answer plus nine revisions). The single correctness decision stops at probability >=0.8; cap3 is a matched secondary prefix. All planned questions and failures remain in denominators; all 20 completed with no missing episodes.

Jev model jev-1.13.0 receives the same complete public question, evidence and current draft; reference answers remain evaluator-only. Its provider 64K total context differs from local Strands 4,096 tokens. No inputs were truncated, no head weights or temperatures were refitted, and no new writer or local encoder work occurred. Seed42 applies to cohort/writer/bootstrap; hosted requests expose no deterministic seed guarantee.

Policy	Complete	EM	F1	Drafts/q	Repairs	Regressions
Fixed1	20/20	50%	0.540	1.00	0	0
Fixed2	20/20	50%	0.540	2.00	0	0
Fixed3	20/20	50%	0.540	3.00	0	0
Fixed5	20/20	50%	0.540	5.00	0	0
Fixed10	20/20	50%	0.540	10.00	0	0
Strands released, cap10	20/20	50%	0.540	5.50	0	0
Strands trained, cap10	20/20	50%	0.540	5.50	0	0
Strands released, cap3	20/20	50%	0.540	2.00	0	0
Strands trained, cap3	20/20	50%	0.540	2.00	0	0
Jev 1.13, cap10	20/20	50%	0.540	3.75	0	0
Jev 1.13, cap3	20/20	50%	0.540	1.65	0	0

Approval precision uses a different denominator

First-decision precision describes accepted first drafts. Final stopping precision includes later approvals. Neither is all-question answer EM. Jev first approves 13 drafts (10 correct); its final approval set has 14 drafts (10 correct). Strands approves the same 10 first drafts in each head, with five correct. Read every conditional precision beside its approval coverage.

Policy / scope	Approved/N	Correct approved	Wrong approved	Precision	Coverage
Jev: FIRST decision	13/20	10	3	76.92%	65%
Jev: FINAL stopping decision	14/20	10	4	71.43%	70%
Strands released: first and final	10/20	5	5	50.00%	50%
Strands trained: first and final	10/20	5	5	50.00%	50%

First-decision probability diagnostics

All20 first decisions were observed in every arm. AUROC ranks EM correctness; higher is better. Brier and NLL assess probabilities; lower is better. Jev supplies provider Noul probabilities, with raw logits and temperature null. Its probability-only NLL clips p to [1e-15,1-1e-15]; 0/20 endpoint values required clipping. Strands NLL uses recorded logits and each saved CAL temperature. No hosted calibration or fabricated logits are claimed.

The journal retains validated adapter probabilities and usage, rather than original HTTP response bodies. Better observed correctness judging did not supply a correct replacement answer for the writer errors. Approval precision and model probability are distinct from answer accuracy.

Policy	Scored/N	AUROC	Brier	NLL
Jev 1.13, cap10	20/20	0.915000	0.141200	0.469715
Strands released, cap10	20/20	0.380000	0.336260	0.910647
Strands trained, cap10	20/20	0.380000	0.336322	0.910822

Secondary paired probability differences

Jev minus each calibrated Strands head; original question is the pairing unit, 5,000 percentile-bootstrap draws, seed42, all20 observed pairs. Precision and coverage below refer to FIRST decisions. These are descriptive secondary intervals without multiplicity correction; all5,000 AUROC/precision draws were defined (zero undefined). They do not establish broad generalization or improved answer quality.

Jev minus Strands	First-decision statistic	Delta	Paired95% interval
released	auroc	+0.535000	[+0.275000, +0.770881]
released	brier	-0.195060	[-0.329511, -0.064655]
released	nll	-0.440933	[-0.802569, -0.075205]
released	approval_precision	+0.269231	[-0.047619, +0.583333]
released	approval_coverage	+0.150000	[-0.200000, +0.450000]
trained	auroc	+0.535000	[+0.275000, +0.770881]
trained	brier	-0.195122	[-0.329620, -0.064704]

Jev minus Strands	First-decision statistic	Delta	Paired95% interval
trained	nll	-0.441107	[-0.802758, -0.075290]
trained	approval_precision	+0.269231	[-0.047619, +0.583333]
trained	approval_coverage	+0.150000	[-0.200000, +0.450000]

Paired answer-quality and draft-use differences

Jev minus comparison; EM/F1 deltas and their intervals are percentage points. Draft deltas are drafts/question. All20 questions are retained, with missing/failed answer quality scored zero. Every EM/F1 delta is zero with [0,0] because per-question scores coincide; this degeneracy does not prove equivalence. Bootstrap: 5,000 draws, seed42, original question as unit.

Jev uses 3.75 drafts/question versus 5.50 for either Strands head; at cap3 the means are 1.65 versus 2.00. The draft-use intervals include zero. Jev still uses more drafts than Fixed1 without improving EM/F1. Token and prefix-model-time comparisons are replay counterfactuals, not measured live latency, deployment savings or equivalent compute across tokenizers/hardware.

Contrast	EM pp [95%CI]	F1 pp [95%CI]	Draft delta [95%CI]
Jev cap10 minus Fixed1	0.00 [0.00, 0.00]	0.00 [0.00, 0.00]	+2.75 [+1.00, +4.55]
Jev cap10 minus released	0.00 [0.00, 0.00]	0.00 [0.00, 0.00]	-1.75 [-4.85, +1.35]
Jev cap10 minus trained	0.00 [0.00, 0.00]	0.00 [0.00, 0.00]	-1.75 [-4.85, +1.35]
Jev cap3 minus released cap3	0.00 [0.00, 0.00]	0.00 [0.00, 0.00]	-0.35 [-1.00, +0.30]
Jev cap3 minus trained cap3	0.00 [0.00, 0.00]	0.00 [0.00, 0.00]	-0.35 [-1.00, +0.30]

Incremental hosted expense and elapsed scopes

The optional phase made 75 unique successful hosted calls: 163,044 input tokens, 1,500 output tokens and recorded charge $0.006847848. No pending/error/uncertain request, unknown usage, duplicate dispatch or retry remained. Finalized hosted accounting and exact ordered journal-to-ledger token/nanodollar matches were independently verified.

The existing cumulative ledger moved from 531 settled calls / $0.047882184 to 606 / $0.054730032: delta 75 calls / $0.006847848, zero reserved/unsettled. The same ledger preserves its $4.50 soft limit under the $5 total authorization; it was never reset. Charges are project accounting at the recorded rate, not a total system-cost or invoice estimate.

Hosted physical work is counted once. Cap3 and fixed controls reuse the same calls/trajectories; the optional phase adds ZERO new writer requests and ZERO local encoder forwards. Primary writer/encoding/training expenses remain in the separate primary report. Local hardware and electricity remain unpriced.

Timing scope	Seconds	Interpretation
Optional replay phase wall	392.261830	Includes work, rests and process overhead within this phase
Intentional idle	375.016000	Inside phase wall; not added again

Timing scope	Seconds	Interpretation
Journal request time	13.184470	Inside phase wall; contains provider roundtrip
Provider roundtrip time	12.217330	Inside request and phase wall; not added again

Concrete failure and limits

For the Gilgamesh question about the divinity associated with fermentation, the reference is "the Waters of Death" and both first/tenth writer answers are "the shore". Jev assigns 0.88 and approves round1; released/trained Strands assign approximately 0.599066/0.599091 and continue to round10. All final answers remain EM-wrong. A correctness critic cannot repair a missing correct candidate through stopping alone.

This is a 20-question public-validation panel and an optional observational comparison after the predeclared primary protocol. Repeated drafts do not increase independent N. EM is a normalized answer-agreement proxy rather than independent support/completeness annotation. Larger disjoint panels and matched live feedback/routing experiments are required before deployment, repair or generalized superiority claims. The primary canonical comparison and released/trained heads remain unchanged.

Evidence references

Input SHA256s and missing-file records are preserved in the provenance JSON. Source paths below identify the actual receipts; they are not model outputs.

[R1] Independent optional audit: OPTIONAL-JEV-FRESH20-INDEPENDENT-TERMINAL-AUDIT-20261018.json (SHA256 809c1f6c54b7...)

[R2] Readable audited result: OPTIONAL-JEV-FRESH20-INDEPENDENT-RESULTS-20261018.txt (SHA256 5c99bf48e595...)

[R3] Separate primary professor report: Primary terminal report, 20:46 Central (SHA256 38caa18fde31...)

[R4] Independent primary audit: FRESH20-INDEPENDENT-TERMINAL-AUDIT-20261018.json (SHA256 c7c590e0629c...)